

# Tanzanian Water Wells

## Problem Description



Water Aid is an NGO based in the United Kingdom that works on access to clean water around the world. They consider access to clean water, decent toilets and good hygiene as basic human rights. For over 30 years, they have been working in partnership to improve access to these three essentials through a combination of programmatic and policy work.

Water Aid works in several countries around the globe, including Tanzania. According to the World Sector Report (2019) around 60% of Tanzanians have access to improved water, but the degree of water access, and the water quality and quantity, varies. Drought, landscape change, and the amplifying effects of climate change are straining existing surface water supplies.

Water Aid is launching a program to repair non-functioning wells in the cross country shared water basins of Eastern Africa. The status of the wells is not clearly recorded in countries surrounding Tanzania. Identifying non-functioning wells, securing funding, and traveling to these rural locations to repair wells is both time and resource intensive. They need a predictive model that accurately identifies which wells are not functioning to reduce cost and ensure they are using their resources wisely. They also need to identify a specific water basin to begin their work.

## Goals

There are three data science goals to address Water Aid's need for accurately identifying non-functioning wells:

1. Using an iterative process, build a predictive machine learning model based on existing water well data to accurately classify non-functioning wells.
2. Deliver two recommendations to Water Aid: a specific transboundary water basin to begin their operations, and one feature characteristic of the wells in this basin that will

lead to higher chance of identifying non-functioning wells.

### Load Packages and Data

```
In [426]: 1 import pandas as pd
          2 import numpy as np
          3 from matplotlib import pyplot as plt
          4 import seaborn as sns
          5
          6 %matplotlib inline
          7
          8 from sklearn.model_selection import train_test_split, GridSearchCV
          9 from sklearn.pipeline import Pipeline
         10 from sklearn.preprocessing import StandardScaler, OneHotEncoder, F
         11 from sklearn.impute import SimpleImputer
         12 from sklearn.compose import ColumnTransformer
         13 from sklearn.linear_model import LogisticRegression
         14 from sklearn.tree import DecisionTreeClassifier
         15 from sklearn.ensemble import RandomForestClassifier, GradientBoost
         16 from sklearn.metrics import plot_confusion_matrix, recall_score,\
         17     accuracy_score, precision_score, f1_score
         18
         19 from imblearn.over_sampling import SMOTE
         20 from imblearn.pipeline import Pipeline as ImPipeline
```

```
In [427]: 1 # Load the predictor data
          2
          3 wells = pd.read_csv('training_set_values.csv')
          4 wells.head()
```

```
Out[427]:
```

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	w
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825359	

5 rows × 10 columns

```
In [428]: 1 # Load the target data
          2
          3 target = pd.read_csv('training_set_labels.csv')
          4
          5 target.head()
```

Out[428]:

	id	status_group
0	69572	functional
1	8776	functional
2	34310	functional
3	67743	non functional
4	19728	functional

## 1. Exploratory Data Analysis

Get a sense of the big picture for the dataset. Prepare the data for further analysis. Gain an understanding of the variables, or predictors in this case. Study the relationship between variables. Make plan for initial model.

```
In [429]: 1 # Identify size of dataset
          2
          3 print("Records for wells:", wells.shape)
          4 print()
          5 print("Records for target:", target.shape)
```

Records for wells: (59400, 40)

Records for target: (59400, 2)

```
In [430]: 1 # Identify datatypes and record amount for each predictor
          2
          3 wells.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59400 entries, 0 to 59399
Data columns (total 40 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     59400 non-null  int64
1   amount_tsh                           59400 non-null  float64
2   date_recorded                        59400 non-null  object
3   funder                               55765 non-null  object
4   gps_height                           59400 non-null  int64
5   installer                           55745 non-null  object
6   longitude                            59400 non-null  float64
7   latitude                             59400 non-null  float64
8   wpt_name                             59400 non-null  object
9   num_private                          59400 non-null  int64
10  basin                                59400 non-null  object
11  subvillage                          59029 non-null  object
12  region                              59400 non-null  object
13  region_code                         59400 non-null  int64
14  district_code                      59400 non-null  int64
15  lga                                 59400 non-null  object
16  ward                                59400 non-null  object
17  population                          59400 non-null  int64
18  public_meeting                     56066 non-null  object
19  recorded_by                        59400 non-null  object
20  scheme_management                  55523 non-null  object
21  scheme_name                       31234 non-null  object
22  permit                            56344 non-null  object
23  construction_year                 59400 non-null  int64
24  extraction_type                   59400 non-null  object
25  extraction_type_group             59400 non-null  object
26  extraction_type_class             59400 non-null  object
27  management                        59400 non-null  object
28  management_group                 59400 non-null  object
29  payment                           59400 non-null  object
30  payment_type                     59400 non-null  object
31  water_quality                     59400 non-null  object
32  quality_group                     59400 non-null  object
33  quantity                          59400 non-null  object
34  quantity_group                   59400 non-null  object
35  source                            59400 non-null  object
36  source_type                       59400 non-null  object
37  source_class                      59400 non-null  object
38  waterpoint_type                   59400 non-null  object
39  waterpoint_type_group            59400 non-null  object
dtypes: float64(3), int64(7), object(30)
memory usage: 18.1+ MB
```

```
In [431]: 1 # Examine numerical predictors mean, min, max
          2
          3 wells.describe()
```

Out[431]:

	id	amount_tsh	gps_height	longitude	latitude	num_private
<b>count</b>	59400.000000	59400.000000	59400.000000	59400.000000	5.940000e+04	59400.000000
<b>mean</b>	37115.131768	317.650385	668.297239	34.077427	-5.706033e+00	0.474141
<b>std</b>	21453.128371	2997.574558	693.116350	6.567432	2.946019e+00	12.236230
<b>min</b>	0.000000	0.000000	-90.000000	0.000000	-1.164944e+01	0.000000
<b>25%</b>	18519.750000	0.000000	0.000000	33.090347	-8.540621e+00	0.000000
<b>50%</b>	37061.500000	0.000000	369.000000	34.908743	-5.021597e+00	0.000000
<b>75%</b>	55656.500000	20.000000	1319.250000	37.178387	-3.326156e+00	0.000000
<b>max</b>	74247.000000	350000.000000	2770.000000	40.345193	-2.000000e-08	1776.000000

```
In [432]: 1 # Missing data total
          2 wells.isna().sum().sum()
```

Out[432]: 46094

```
In [433]: 1 # Missing data by predictor
          2
          3 wells.isna().sum()
```

```
Out[433]: id                                0
          amount_tsh                        0
          date_recorded                     0
          funder                            3635
          gps_height                         0
          installer                         3655
          longitude                         0
          latitude                          0
          wpt_name                          0
          num_private                        0
          basin                             0
          subvillage                        371
          region                            0
          region_code                       0
          district_code                     0
          lga                               0
          ward                              0
          population                        0
          public_meeting                    3334
          recorded_by                       0
          scheme_management                  3877
          scheme_name                       28166
          permit                            3056
          construction_year                  0
          extraction_type                    0
          extraction_type_group              0
          extraction_type_class              0
          management                         0
          management_group                   0
          payment                           0
          payment_type                       0
          water_quality                      0
          quality_group                      0
          quantity                          0
          quantity_group                     0
          source                             0
          source_type                        0
          source_class                       0
          waterpoint_type                    0
          waterpoint_type_group              0
          dtype: int64
```

```
In [434]: 1 # Data missing for target
          2 target.isna().sum()
```

```
Out[434]: id                                0
          status_group                       0
          dtype: int64
```


```
In [435]: 1 # Examine value counts for the target, consider imbalance in target
          2 target['status_group'].value_counts()
```

```
Out[435]: functional                32259
          non functional            22824
          functional needs repair    4317
          Name: status_group, dtype: int64
```

```
In [436]: 1 # Percentage makeup of target values
          2 print("Functional percentage:", round(32259/59400*100, 2))
          3 print("Non functional percentage:", round(22824/59400*100, 2))
          4 print("Functional needs repair percentage:", round(4317/59400*100,
```

```
Functional percentage: 54.31
Non functional percentage: 38.42
Functional needs repair percentage: 7.27
```

```
In [437]: 1 # Visually plot target variable counts
          2
          3 target.status_group.value_counts().plot(kind="bar")
          4 plt.title("Number of Functioning Wells", fontsize= 18)
          5 plt.xlabel("Well Functionality", fontsize = 12)
          6 plt.xticks(rotation=0)
          7 plt.ylabel("Count", fontsize = 12)
          8 plt.show();
          9
         10 plt.savefig('Number of Functioning Wells')
```



<Figure size 432x288 with 0 Axes>

```
In [438]: 1 # Identify unique values per column  
          2 print(wells.nunique())
```

```
id                59400  
amount_tsh        98  
date_recorded     356  
funder            1897  
gps_height        2428  
installer         2145  
longitude         57516  
latitude          57517  
wpt_name          37400  
num_private       65  
basin             9  
subvillage        19287  
region            21  
region_code       27  
district_code     20  
lga               125  
ward              2092  
population        1049  
public_meeting    2  
recorded_by       1  
scheme_management 12  
scheme_name       2696  
permit            2  
construction_year 55  
extraction_type    18  
extraction_type_group 13  
extraction_type_class 7  
management         12  
management_group   5  
payment            7  
payment_type       7  
water_quality      8  
quality_group      6  
quantity           5  
quantity_group     5  
source             10  
source_type        7  
source_class       3  
waterpoint_type    7  
waterpoint_type_group 6  
dtype: int64
```



```
In [439]: 1 # Concatenate preds and target for heatmap
          2
          3 df = pd.concat([wells, target], axis =1)
          4
          5 df = df.loc[:,~df.columns.duplicated()].copy()
          6
          7 df.head()
```

Out [439]:

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	v
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825359	


5 rows × 41 columns

### Correlation of numeric data

```

In [440]: 1 # Create a heatmap to examine the correlational coefficients
          2
          3
          4 corr = df.corr()
          5
          6 # Set up figure and axes
          7 fig, ax = plt.subplots(figsize=(8, 12))
          8
          9 # Plot a heatmap of the correlations
         10
         11 sns.heatmap(
         12     data=corr,
         13     mask=np.triu(np.ones_like(corr, dtype=bool)),
         14     ax=ax,
         15     annot=True,
         16     # Customizes colorbar appearance
         17     cbar_kws={"label": "Correlation", "orientation": "horizontal",
         18             }
         19 )
         20
         21 # Customize the plot appearance
         22
         23
         24 ax.set_title("Heatmap of Correlation Between Attributes (Including
         25

```




```

In [441]: 1 # Check distribution for numeric data
          2 import warnings
          3 warnings.filterwarnings('ignore')
          4
          5 print(sns.distplot(wells.population, bins = 100))
          6
          7

```

AxesSubplot(0.125,0.125;0.775x0.755)




```

In [442]: 1 # Identify range of population and any outliers
          2
          3 sns.boxplot(wells.population)
          4

```

Out[442]: <AxesSubplot:xlabel='population'>



```
In [443]: 1 # Identify population counts
          2 print(wells.population.value_counts())
          3 print(wells.population.nunique())
```

0	21381
1	7025
200	1940
150	1892
250	1681
...	
3241	1
1960	1
1685	1
2248	1
1439	1


Name: population, Length: 1049, dtype: int64  
1049

### Total Static Head data

```
In [444]: 1 # Plot and describe total static head
          2
          3 sns.distplot(wells.amount_tsh)
          4
          5 print(wells.amount_tsh.describe())
          6
```


count	59400.000000
mean	317.650385
std	2997.574558
min	0.000000
25%	0.000000
50%	0.000000
75%	20.000000
max	350000.000000

Name: amount\_tsh, dtype: float64



```
In [445]: 1 # Identify range and outliers of total static head
          2
          3 sns.boxplot(wells.amount_tsh)
```

```
Out[445]: <AxesSubplot:xlabel='amount_tsh'>
```




```
In [446]: 1 # Identify how many wells do not have static head
          2
          3 print(wells.amount_tsh.value_counts())
```

```
0.0          41639
500.0         3102
50.0          2472
1000.0        1488
20.0          1463
...
8500.0         1
6300.0         1
220.0          1
138000.0       1
12.0           1
Name: amount_tsh, Length: 98, dtype: int64
```

### Geographic Data

```
In [447]: 1 # plot basins
2 fig, axs = plt.subplots(1, 1,
3                       figsize=(6, 4),
4                       tight_layout = True)
5
6 axs.hist(wells.basin, bins = 18)
7 plt.xticks(rotation = 90)
8
9 plt.show()
```



```
In [448]: 1 # How many wells are in each basin
2
3 wells.basin.value_counts()
```

```
Out[448]: Lake Victoria      10248
Pangani      8940
Rufiji      7976
Internal     7785
Lake Tanganyika 6432
Wami / Ruvu  5987
Lake Nyasa   5085
Ruvuma / Southern Coast 4493
Lake Rukwa   2454
Name: basin, dtype: int64
```

```
In [449]: 1  
          2 wells.region.value_counts()
```

```
Out[449]: Iringa          5294  
          Shinyanga      4982  
          Mbeya          4639  
          Kilimanjaro    4379  
          Morogoro       4006  
          Arusha         3350  
          Kagera         3316  
          Mwanza         3102  
          Kigoma         2816  
          Ruvuma         2640  
          Pwani          2635  
          Tanga          2547  
          Dodoma         2201  
          Singida        2093  
          Mara           1969  
          Tabora         1959  
          Rukwa           1808  
          Mtwara         1730  
          Manyara       1583  
          Lindi          1546  
          Dar es Salaam   805  
          Name: region, dtype: int64
```

```
In [450]: 1 wells.district_code.value_counts()
```

```
Out[450]: 1      12203  
          2      11173  
          3       9998  
          4       8999  
          5       4356  
          6       4074  
          7       3343  
          8       1043  
          30       995  
          33       874  
          53       745  
          43       505  
          13       391  
          23       293  
          63       195  
          62       109  
          60        63  
          0         23  
          80        12  
          67         6  
          Name: district_code, dtype: int64
```

```
In [451]: 1 wells.region_code.value_counts()
```

```
Out[451]: 11    5300
          17    5011
          12    4639
           3    4379
           5    4040
          18    3324
          19    3047
           2    3024
          16    2816
          10    2640
           4    2513
           1    2201
          13    2093
          14    1979
          20    1969
          15    1808
           6    1609
          21    1583
          80    1238
          60    1025
          90     917
           7     805
          99     423
           9     390
          24     326
           8     300
          40        1
          Name: region_code, dtype: int64
```

### Water attributes

```
In [452]: 1 wells.water_quality.value_counts()
```

```
Out[452]: soft                50818
          salty                4856
          unknown             1876
          milky                804
          coloured             490
          salty abandoned      339
          fluoride             200
          fluoride abandoned    17
          Name: water_quality, dtype: int64
```



```
In [453]: 1 wells.quality_group.value_counts()
```

```
Out[453]: good          50818
          salty         5195
          unknown      1876
          milky         804
          colored       490
          fluoride      217
          Name: quality_group, dtype: int64
```

```
In [454]: 1 wells.quantity.value_counts()
```

```
Out[454]: enough        33186
          insufficient   15129
          dry            6246
          seasonal       4050
          unknown        789
          Name: quantity, dtype: int64
```

```
In [455]: 1 wells.quantity_group.value_counts()
```

```
Out[455]: enough        33186
          insufficient   15129
          dry            6246
          seasonal       4050
          unknown        789
          Name: quantity_group, dtype: int64
```

```
In [456]: 1 wells.scheme_name.value_counts()
```

```
Out[456]: K                      682
          None                   644
          Borehole               546
          Chalinze wate         405
          M                      400
          ...
          Mws                    1
          Mpal                   1
          Malemeo gravity water supply 1
          Bulenya water supply      1
          UNICRF                  1
          Name: scheme_name, Length: 2696, dtype: int64
```

In [457]: 1 wells.scheme\_management.value\_counts()

```
Out[457]: VWC          36793
          WUG          5206
          Water authority  3153
          WUA          2883
          Water Board    2748
          Parastatal     1680
          Private operator 1063
          Company        1061
          Other          766
          SWC           97
          Trust          72
          None           1
          Name: scheme_management, dtype: int64
```

In [458]: 1 wells.extraction\_type.value\_counts()

```
Out[458]: gravity          26780
          nira/tanira      8154
          other            6430
          submersible      4764
          swn 80           3670
          mono             2865
          india mark ii    2400
          afridev          1770
          ksb              1415
          other - rope pump  451
          other - swn 81    229
          windmill         117
          india mark iii    98
          cemo              90
          other - play pump  85
          walimi            48
          climax           32
          other - mkulima/shinyanga 2
          Name: extraction_type, dtype: int64
```

In [459]: 1 wells.extraction\_type\_group.value\_counts()

```
Out[459]: gravity          26780
          nira/tanira      8154
          other            6430
          submersible      6179
          swn 80           3670
          mono             2865
          india mark ii    2400
          afridev          1770
          rope pump         451
          other handpump    364
          other motorpump   122
          wind-powered      117
          india mark iii    98
          Name: extraction_type_group, dtype: int64
```

```
In [460]: 1 wells.extraction_type_class.value_counts()
```

```
Out[460]: gravity          26780
handpump          16456
other             6430
submersible       6179
motorpump         2987
rope pump         451
wind-powered      117
Name: extraction_type_class, dtype: int64
```

```
In [461]: 1 wells.source.value_counts()
```

```
Out[461]: spring          17021
shallow well          16824
machine dbh           11075
river                 9612
rainwater harvesting   2295
hand dtw              874
lake                  765
dam                   656
other                 212
unknown               66
Name: source, dtype: int64
```

```
In [462]: 1 wells.source_type.value_counts()
```

```
Out[462]: spring          17021
shallow well          16824
borehole             11949
river/lake           10377
rainwater harvesting   2295
dam                   656
other                 278
Name: source_type, dtype: int64
```

```
In [463]: 1 wells.waterpoint_type_group.value_counts()
```

```
Out[463]: communal standpipe  34625
hand pump                    17488
other                        6380
improved spring              784
cattle trough                116
dam                           7
Name: waterpoint_type_group, dtype: int64
```

```
In [464]: 1 wells.source_class.value_counts()
```

```
Out[464]: groundwater      45794
surface                   13328
unknown                   278
Name: source_class, dtype: int64
```

In [465]: 1 wells.waterpoint\_type.value\_counts()

```
Out[465]: communal standpipe      28522
hand pump      17488
other          6380
communal standpipe multiple    6103
improved spring      784
cattle trough      116
dam                7
Name: waterpoint_type, dtype: int64
```

### Organizational attributes

In [466]: 1 wells.funder.value\_counts()

```
Out[466]: Government Of Tanzania      9084
Danida      3114
Hesawa      2202
Rwssp       1374
World Bank  1349
...
Fida        1
Kigoma Municipal Council      1
Abc-ihushi Development Cent    1
Tag Church Ub      1
Nyamingu Subvillage      1
Name: funder, Length: 1897, dtype: int64
```

In [467]: 1 wells.num\_private.value\_counts()

```
Out[467]: 0      58643
6         81
1         73
5         46
8         46
...
180        1
213        1
23         1
55         1
94         1
Name: num_private, Length: 65, dtype: int64
```

In [468]: 1 wells.permit.value\_counts()

```
Out[468]: True      38852
False     17492
Name: permit, dtype: int64
```

```
In [469]: 1 wells.management.value_counts()
```

```
Out[469]: vwc          40507
          wug           6515
          water board   2933
          wua           2535
          private operator 1971
          parastatal    1768
          water authority  904
          other         844
          company       685
          unknown       561
          other - school  99
          trust         78
          Name: management, dtype: int64
```

```
In [470]: 1 wells.management_group.value_counts()
```

```
Out[470]: user-group    52490
          commercial    3638
          parastatal    1768
          other         943
          unknown       561
          Name: management_group, dtype: int64
```

```
In [471]: 1 wells.payment.value_counts()
```

```
Out[471]: never pay      25348
          pay per bucket  8985
          pay monthly     8300
          unknown         8157
          pay when scheme fails 3914
          pay annually     3642
          other           1054
          Name: payment, dtype: int64
```

```
In [472]: 1 wells.payment_type.value_counts()
```

```
Out[472]: never pay      25348
          per bucket     8985
          monthly        8300
          unknown         8157
          on failure      3914
          annually        3642
          other           1054
          Name: payment_type, dtype: int64
```

```
In [473]: 1 with pd.option_context('display.max_rows', 5, 'display.max_columns'
2          display(wells[1000:1020]))
```

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	lati
1000	47384	250.0	2013-02-14	Oxfam	1409	OXFAM	30.105401	-4.367
1001	11570	0.0	2012-10-12	Resolute Mining	0	Consulting Engineer	33.210098	-4.049
...	...	...	...	...	...	...	...	...
1018	41433	0.0	2011-03-05	Government Of Tanzania	1307	DWE	38.325050	-4.464
1019	21810	0.0	2013-01-17	Bulyahunlu Gold Mine	0	Bulyahunlu Gold Mine	32.370100	-3.287

20 rows × 40 columns

## 2. Preprocess data, Initial Model

Redundant data columns where the data is included in other columns that contain more expansive information should be dropped: water attributes, geographic attributes, include water include regional columns, water extraction and source types.

Drop columns that do not contribute to the model. These include water id, names of the waterpoint, names of subvillages.

Make plan for missing categorical and numeric data.

```
In [474]: 1 # Drop redundant data columns and columns that do not contribute t
2 wells.drop(columns = ['id', 'wpt_name', 'region', 'recorded_by', '
3                 'scheme_management', 'extraction_type_group', 'payment_
4                 'quality_group', 'quantity_group', 'source_type', 'wate
5
```

In [475]: 1 wells.columns

Out[475]: Index(['amount\_tsh', 'date\_recorded', 'funder', 'gps\_height', 'installer',  
'longitude', 'latitude', 'num\_private', 'basin', 'region\_code',  
'district\_code', 'lga', 'ward', 'population', 'public\_meeting',  
'permit', 'construction\_year', 'extraction\_type',  
'extraction\_type\_class', 'management', 'management\_group', 'payment',  
'water\_quality', 'quantity', 'source', 'source\_class',  
'waterpoint\_type\_group'],  
dtype='object')

In [476]: 1 *# Check missing data*  
2 wells.isna().sum()

Out[476]: amount\_tsh 0  
date\_recorded 0  
funder 3635  
gps\_height 0  
installer 3655  
longitude 0  
latitude 0  
num\_private 0  
basin 0  
region\_code 0  
district\_code 0  
lga 0  
ward 0  
population 0  
public\_meeting 3334  
permit 3056  
construction\_year 0  
extraction\_type 0  
extraction\_type\_class 0  
management 0  
management\_group 0  
payment 0  
water\_quality 0  
quantity 0  
source 0  
source\_class 0  
waterpoint\_type\_group 0  
dtype: int64

```
In [477]: 1 # Replace Nan in public_meeting and permit as False
          2
          3 wells['public_meeting'] = wells['public_meeting'].fillna('False').
          4 wells.public_meeting.head()
```

```
Out[477]: 0    True
          1    True
          2    True
          3    True
          4    True
          Name: public_meeting, dtype: bool
```

```
In [478]: 1 # replace missing permit data as False
          2 wells['permit'] = wells['permit'].fillna('False').astype('bool')
          3 wells.permit.head()
```

```
Out[478]: 0    False
          1     True
          2     True
          3     True
          4     True
          Name: permit, dtype: bool
```

```
In [479]: 1 # Convert "date_recorded" to month_recorded
          2
          3 import datetime
          4
          5 wells['date_recorded'] = pd.to_datetime(wells['date_recorded'])
          6 wells['month_recorded'] = wells['date_recorded'].dt.month
          7 wells['month_recorded']
```

```
Out[479]: 0         3
          1         3
          2         2
          3         1
          4         7
          ..
          59395     5
          59396     5
          59397     4
          59398     3
          59399     3
          Name: month_recorded, Length: 59400, dtype: int64
```

```
In [480]: 1 wells.drop('date_recorded', axis = 1, inplace = True)
```

### Initial Model - Logistic Regression

Use a Logistic Regression model in a pipeline for initial model results.



```
In [481]: 1 # Assign the predictors and target
          2 X = wells
          3 y = target['status_group']
```

```
In [482]: 1 # Perform a train test split
          2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [483]: 1 X_train.columns
```

```
Out[483]: Index(['amount_tsh', 'funder', 'gps_height', 'installer', 'longitude',
                'latitude', 'num_private', 'basin', 'region_code', 'district_code',
                'lga', 'ward', 'population', 'public_meeting', 'permit',
                'construction_year', 'extraction_type', 'extraction_type_class',
                'management', 'management_group', 'payment', 'water_quality',
                'quantity', 'source', 'source_class', 'waterpoint_type_group',
                'month_recorded'],
                dtype='object')
```

```
In [484]: 1 # Examine data types and record counts
          2 X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 44550 entries, 24947 to 56422
Data columns (total 27 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   amount_tsh                             44550 non-null   float64
 1   funder                                 41859 non-null   object
 2   gps_height                             44550 non-null   int64
 3   installer                             41850 non-null   object
 4   longitude                             44550 non-null   float64
 5   latitude                              44550 non-null   float64
 6   num_private                           44550 non-null   int64
 7   basin                                 44550 non-null   object
 8   region_code                           44550 non-null   int64
 9   district_code                         44550 non-null   int64
10   lga                                   44550 non-null   object
11   ward                                 44550 non-null   object
12   population                           44550 non-null   int64
13   public_meeting                       44550 non-null   bool
14   permit                               44550 non-null   bool
15   construction_year                   44550 non-null   int64
16   extraction_type                     44550 non-null   object
17   extraction_type_class                44550 non-null   object
18   management                           44550 non-null   object
19   management_group                     44550 non-null   object
20   payment                             44550 non-null   object
21   water_quality                       44550 non-null   object
22   quantity                             44550 non-null   object
23   source                               44550 non-null   object
24   source_class                         44550 non-null   object
25   waterpoint_type_group                44550 non-null   object
26   month_recorded                       44550 non-null   int64
dtypes: bool(2), float64(3), int64(7), object(15)
memory usage: 8.9+ MB
```

In [539]:

```
1
2
3 # create subpipe for numeric data
4
5 subpipe_num = Pipeline(steps=[('num_impute', SimpleImputer()),
6                               ('ss', StandardScaler())])
7
8 # create subpipe for categorical data, use SimpleImputer for 'miss
9
10 subpipe_cat = Pipeline(steps=[('cat_impute', SimpleImputer(strateg
11                           ('ohe', OneHotEncoder(sparse=False, h
12
13 # combine subpipes into ColumnTransformer
14
15 CT = ColumnTransformer(transformers=[('subpipe_num', subpipe_num,
16                                     ('subpipe_cat', subpipe_cat, [
17
18
19
20                               remainder='passthrough')
21
22
```

In [540]:

```
1 #Perform Logistic Regression for initial model
2
3 log_reg_pipe = Pipeline(steps = [('ct', CT),
4                                  ('log_reg', LogisticRegression(random_s
```

```
In [541]: 1 # Fit the logistic regression model
          2 log_reg_pipe.fit(X_train, y_train)
```

```
Out[541]: Pipeline(steps=[('ct',
                             ColumnTransformer(remainder='passthrough',
                                                    transformers=[('subpipe_num',
                                                                    Pipeline(steps=[('n
um_impute',
                                                                    Si
mpleImputer()),
                                                                    ('s
s',
                                                                    St
andardScaler())])),
                             [0, 2, 4, 5, 12]),
                            ('subpipe_cat',
                             Pipeline(steps=[('c
at_impute',
                                                                    Si
mpleImputer(fill_value='missing',
                                                                    strategy='constant'))],
                                                                    ('o
he',
                                                                    On
eHotEncoder(handle_unknown='ignore',
                                                                    sparse=False))])),
                             [1, 3, 6, 7, 8, 9,
                             10, 11, 13,
                             14, 15, 16, 17, 1
                             8, 19, 20,
                             21, 22, 23, 24, 2
                             5, 26])])),
                ('log_reg', LogisticRegression(random_state=42)))]
```

### Evaluate initial model

```
In [542]: 1 # Score the log reg model
          2 log_reg_pipe.score(X_train, y_train)
```

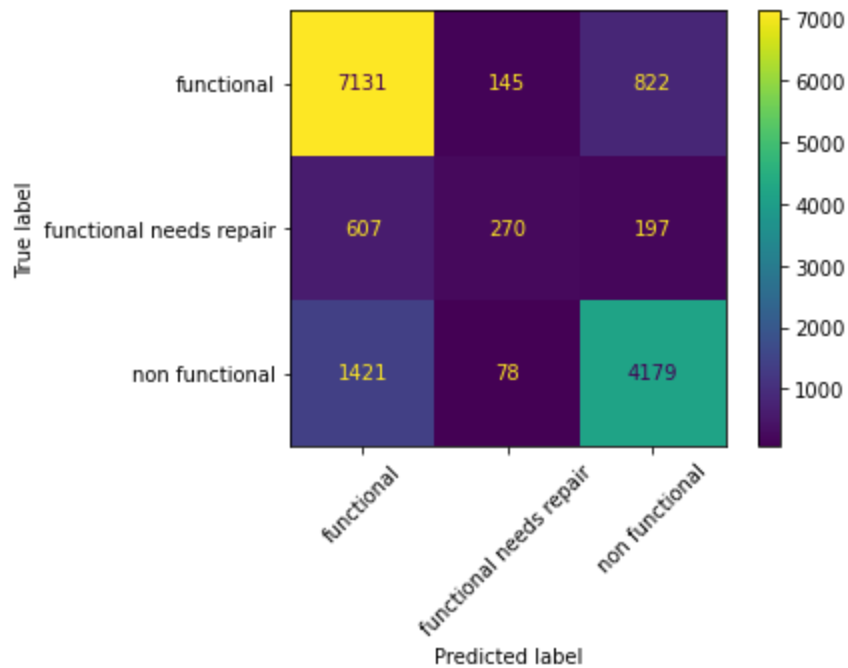
```
Out[542]: 0.8025813692480359
```

```
In [543]: 1 # create predicted target variable
          2 y_hat = log_reg_pipe.predict(X_test)
```

```
In [544]: 1 # Generate log_reg classification report
          2 from sklearn.metrics import classification_report
          3
          4 print(classification_report(y_test, y_hat))
```

	precision	recall	f1-score	support
functional	0.78	0.88	0.83	8098
functional needs repair	0.55	0.25	0.34	1074
non functional	0.80	0.74	0.77	5678
accuracy			0.78	14850
macro avg	0.71	0.62	0.65	14850
weighted avg	0.77	0.78	0.77	14850

```
In [545]: 1 plot_confusion_matrix(log_reg_pipe, X_test, y_test, xticks_rotatio
```



```
In [546]: 1 # Save model in Joblib
2 from joblib import Parallel, delayed
3 import joblib
4
5 import pickle
6
7 # Save the model as a pickle in a file
8 joblib.dump(log_reg_pipe, 'log_reg.pkl')
9
10 # Load the model from the file
11 #log_reg_from_joblib = joblib.load('log_reg.pkl')
12
13 # Use the loaded model to make predictions
14 #log_reg_from_joblib.predict(X_test)
```

```
Out[546]: ['log_reg.pkl']
```

### 3. Decision Tree Model with Parameter Tuning

Considering the dataset a decision tree would be a useful secondary model. Use hyperparameter tuning to improve upon the initial logistic regression model.

```
In [547]: 1 import category_encoders as ce
2
3
4 # create subpipe for numeric data
5
6 subpipe_num = Pipeline(steps=[('num_impute', SimpleImputer()),
7                               ('ss', StandardScaler())])
8
9 # create subpipe for categorical data, use SimpleImputer for 'miss
10
11 subpipe_cat = Pipeline(steps=[('cat_impute', SimpleImputer(strateg
12                               ('ohe', OneHotEncoder(sparse=False, h
13
14 # combine subpipes into ColumnTransformer
15
16 CT = ColumnTransformer(transformers=[('subpipe_num', subpipe_num,
17                                     ('subpipe_cat', subpipe_cat, [
18
19
20
21                                     remainder='passthrough')
22
```

```
In [548]: 1 # Use a decision tree for the secondary model
2 dtc = DecisionTreeClassifier(random_state = 42)
3
4 dtc_pipe = Pipeline(steps=[('ct', CT),
5                             ('dtc', dtc)])
```

In [549]: 1 dtc\_pipe.fit(X\_train, y\_train)

```
Out[549]: Pipeline(steps=[('ct',
                           ColumnTransformer(remainder='passthrough',
                                                transformers=[('subpipe_num',
                                                                Pipeline(steps=[('n
um_impute',
                                                                Si
mpleImputer()),
                                                                ('s
s',
                                                                St
andardScaler())])),
                           [0, 2, 4, 5, 12]),
                           ('subpipe_cat',
                            Pipeline(steps=[('c
at_impute',
                            Si
mpleImputer(fill_value='missing',
strategy='constant'))],
                            ('o
he',
                            On
eHotEncoder(handle_unknown='ignore',
sparse=False))])),
                           [1, 3, 6, 7, 8, 9,
                           10, 11, 13,
                           14, 15, 16, 17, 1
                           8, 19, 20,
                           21, 22, 23, 24, 2
                           5, 26])])),
              ('dtc', DecisionTreeClassifier(random_state=42)))]
```

In [550]: 1 dtc\_pipe.score(X\_train, y\_train)

Out[550]: 0.9984960718294051

In [551]: 1 dtc\_pipe.score(X\_test, y\_test)

Out[551]: 0.7612121212121212


### Evaluate Decision Tree Model

In [552]: 1 y\_hat = dtc\_pipe.predict(X\_test)

```
In [553]: 1 print(classification_report(y_test, y_hat))
```

	precision	recall	f1-score	support
functional	0.81	0.80	0.80	8098
functional needs repair	0.38	0.39	0.39	1074
non functional	0.77	0.77	0.77	5678
accuracy			0.76	14850
macro avg	0.65	0.65	0.65	14850
weighted avg	0.76	0.76	0.76	14850

```
In [554]: 1 plot_confusion_matrix(dtc_pipe, X_test, y_test, xticks_rotation=45)
```



```
In [555]: 1 len(dtc_pipe.named_steps['dtc'].feature_importances_)
```

```
Out[555]: 5974
```

```
In [556]: 1 model_tree = dtc_pipe.named_steps['dtc']
          2 model_tree.feature_importances_
```

```
Out[556]: array([0.02232117, 0.05443548, 0.12866974, ..., 0.00177732, 0.0003740
                2,
                0.00037427])
```

```
In [558]: 1 # Save the decision tree model as a pickle in a file
          2 joblib.dump(dtc_pipe, 'dtc_pipe.pkl')
```

```
Out[558]: ['dtc_pipe.pkl']
```

## Results



The decision tree model's accuracy was less than the logistic regression model and did not improve upon the logistic regression accuracy, though the f1-score for non-functional wells

**Use gridsearch for hyperparameter tuning.**

```
In [75]: 1 params = {}  
2 params['dtc__criterion'] = ['gini', 'entropy']  
3 params['dtc__min_samples_leaf'] = [1, 3, 5, 7, 10]  
4 params['dtc__max_depth'] = [1, 3, 5, 7, 9]  
5 params['dtc__splitter'] = ['best', 'random']  
6  
7 gs = GridSearchCV(estimator=dtc_pipe,  
8                   param_grid=params,  
9                   cv=3)
```

```
In [76]: 1 gs.fit(X_train, y_train)
```

```
Out[76]: GridSearchCV(cv=3,
                      estimator=Pipeline(steps=[('ct',
                                                  ColumnTransformer(remainder='
passthrough',
                                                  transformer
s=[('subpipe_num',
Pipeline(steps=[('num_impute',
SimpleImputer()),
('ss',
StandardScaler())])),
[0, 2,
4, 5,
12]),
('subpipe_cat',
Pipeline(steps=[('cat_impute',
SimpleImputer(fill_value='missing',
strategy='constant'))),
('ohe',
OneHotEncoder(handle_unknown='ignore',
sparse=False))])),
[1, 3,
6, 7,
8, 9,
10,
11,
13,
14,
15,
16,
17,
```

```
18,  
19,  
20,  
21,  
22,  
23,  
24,  
25,  
26]]))),  
                                ('dtc',  
                                DecisionTreeClassifier(random  
_state=42))]),  
                                param_grid={'dtc__criterion': ['gini', 'entropy'],  
                                            'dtc__max_depth': [1, 3, 5, 7, 9],  
                                            'dtc__min_samples_leaf': [1, 3, 5, 7, 10],  
                                            'dtc__splitter': ['best', 'random']})
```

```
In [77]: 1 # Identify the best parameters  
        2 gs.best_params_
```

```
Out[77]: {'dtc__criterion': 'gini',  
          'dtc__max_depth': 9,  
          'dtc__min_samples_leaf': 5,  
          'dtc__splitter': 'best'}
```

```
In [78]: 1 # Examine cross validation results
        2 gs.cv_results_['mean_test_score']
```

```
Out[78]: array([0.6424018 , 0.6424018 , 0.6424018 , 0.6424018 , 0.6424018 ,
                0.6424018 , 0.6424018 , 0.6424018 , 0.6424018 , 0.6424018 ,
                0.69427609, 0.69441077, 0.69427609, 0.69441077, 0.69429854,
                0.69443322, 0.69427609, 0.69445567, 0.69450056, 0.69461279,
                0.71270483, 0.70826038, 0.71261504, 0.70810325, 0.71265993,
                0.70808081, 0.71272727, 0.7081257 , 0.71261504, 0.70808081,
                0.72489338, 0.72282828, 0.72480359, 0.72255892, 0.72455668,
                0.72253648, 0.72453423, 0.72210999, 0.72430976, 0.72190797,
                0.73719416, 0.73476992, 0.73705948, 0.73429854, 0.73748597,
                0.73297419, 0.73643098, 0.73313131, 0.73542088, 0.73236813,
                0.6424018 , 0.6424018 , 0.6424018 , 0.6424018 , 0.6424018 ,
                0.6424018 , 0.6424018 , 0.6424018 , 0.6424018 , 0.6424018 ,
                0.69342312, 0.6935578 , 0.69342312, 0.6935578 , 0.69342312,
                0.6935578 , 0.69342312, 0.69360269, 0.69351291, 0.69362514,
                0.7006734 , 0.69723906, 0.70060606, 0.6973064 , 0.70056117,
                0.6973064 , 0.70042649, 0.69717172, 0.70038159, 0.69710438,
                0.71353535, 0.71380471, 0.71353535, 0.71360269, 0.71384961,
                0.71378227, 0.71367003, 0.7138945 , 0.71335578, 0.71411897,
                0.73158249, 0.7308642 , 0.73167228, 0.7308642 , 0.73156004,
                0.7308193 , 0.73182941, 0.73021324, 0.73113356, 0.73005612])
```


#### Evaluate decision tree gridsearch results

```
In [79]: 1 y_hat = gs.predict(X_test)
```

```
In [80]: 1 print(classification_report(y_test, y_hat))
```

	precision	recall	f1-score	support
functional	0.70	0.93	0.80	8098
functional needs repair	0.58	0.15	0.23	1074
non functional	0.85	0.57	0.68	5678
accuracy			0.74	14850
macro avg	0.71	0.55	0.57	14850
weighted avg	0.75	0.74	0.72	14850

```
In [81]: 1 plot_confusion_matrix(gs, X_test, y_test, xticks_rotation = 45);
```



```
In [559]: 1 # Save the model as a pickle in a file  
2 joblib.dump(gs, 'grid_search_dtc.pkl')
```

```
Out[559]: ['grid_search_dtc.pkl']
```

## Results

While accuracy decreased overall, the precision score on non-functional wells improved from 77% to 85%. This could be a good model if we only focus on precision score for non-functioning wells. Wells that need repair precision score also improved by 20%, this opens a path to possibly identify wells that could soon be non-functioning.

## 4. Random Forest with SMOTE and Tuning

Use a random forest model to further explore whether the precision or recall score on non-functioning wells can be improved. Address class imbalance issues with SMOTE. Further tune the model using search tools for best hyperparameters.

```
In [82]: 1 # Instantiate a Random Forest Classifier
          2
          3 rfc = RandomForestClassifier(random_state=42)
          4
          5 # Instantiate SMOTE for class imbalance
          6
          7 sm = SMOTE(sampling_strategy = 'auto', random_state = 42)
          8
          9 # Create pipeline
         10
         11 rfc_model_pipe = ImPipeline(steps=[('ct', CT),
         12                                     ('sm', sm),
         13                                     ('rfc', rfc)])
         14
```

```
In [83]: 1 rfc_model_pipe.fit(X_train, y_train)
```

```
Out[83]: Pipeline(steps=[('ct',
                           ColumnTransformer(remainder='passthrough',
                                                transformers=[('subpipe_num',
                                                                Pipeline(steps=[('n
um_impute',
                                                                Si
mpleImputer()),
                                                                ('s
s',
                                                                St
andardScaler())])),
                                                                [0, 2, 4, 5, 12]),
('subpipe_cat',
 Pipeline(steps=[('c
at_impute',
                                                                Si
mpleImputer(fill_value='missing',
strategy='constant'))],
                                                                ('o
he',
                                                                On
eHotEncoder(handle_unknown='ignore',
sparse=False))])),
                                                                [1, 3, 6, 7, 8, 9,
10, 11, 13,
                                                                14, 15, 16, 17, 1
8, 19, 20,
                                                                21, 22, 23, 24, 2
5, 26])])),
          ('sm', SMOTE(random_state=42)),
          ('rfc', RandomForestClassifier(random_state=42))])
```

```
In [84]: 1 rfc_model_pipe.score(X_train, y_train)
```

```
Out[84]: 0.9984960718294051
```


**Evaluate results on Random Forest**

```
In [85]: 1 y_hat_rfc = rfc_model_pipe.predict(X_test)
```

```
In [86]: 1 print(classification_report(y_test, y_hat_rfc))
```

	precision	recall	f1-score	support
functional	0.82	0.84	0.83	8098
functional needs repair	0.43	0.43	0.43	1074
non functional	0.82	0.78	0.80	5678
accuracy			0.79	14850
macro avg	0.69	0.69	0.69	14850
weighted avg	0.79	0.79	0.79	14850

```
In [87]: 1 plot_confusion_matrix(rfc_model_pipe, X_test, y_test, xticks_rotat
```



```
In [560]: 1 # Save the random forest model as a pickle in a file
          2 joblib.dump(rfc_model_pipe, 'rfc_model.pkl')
```

```
Out[560]: ['rfc_model.pkl']
```

**Gridsearch for hyperparameter tuning**

```
In [95]: 1 # Grid Search for better model criteria
2
3 params = {'rfc__n_estimators': [10],
4           'rfc__criterion': ['gini'],
5           'rfc__min_samples_leaf': [1, 5, 10],
6           'rfc__max_depth': [1, 5, 9],
7           'rfc__max_features': [9]
8           }
9
10 gs_rfc = GridSearchCV(estimator=rfc_model_pipe,
11                       param_grid=params, n_jobs = -1,
12                       cv=3)
```



```
In [96]: 1 gs_rfc.fit(X_train, y_train)
```

```
Out[96]: GridSearchCV(cv=3,
                      estimator=Pipeline(steps=[('ct',
                                                  ColumnTransformer(remainder='
passthrough',
                                                  transformer
s=[('subpipe_num',
Pipeline(steps=[('num_impute',
SimpleImputer()),
('ss',
StandardScaler())])),
[0, 2,
4, 5,
12]),
('subpipe_cat',
Pipeline(steps=[('cat_impute',
SimpleImputer(fill_value='missing',
strategy='constant'))),
('ohe',
OneHotEncoder(handle_unknown='ignore',
sparse=False))])),
[1, 3,
6, 7,
8, 9,
10,
11,
13,
14,
15,
16,
17,
```

```

18,
19,
20,
21,
22,
23,
24,
25,
26]]))),
                                ('sm', SMOTE(random_state=4
2)),
                                ('rfc',
                                RandomForestClassifier(random
_state=42))]),
                                n_jobs=-1,
                                param_grid={'rfc__criterion': ['gini'],
                                'rfc__max_depth': [1, 5, 9], 'rfc__max_featu
res': [9],
                                'rfc__min_samples_leaf': [1, 5, 10],
                                'rfc__n_estimators': [10]})

```

```

In [103]: 1 # Best parameters for further tuning
          2 gs_rfc.best_params_

```

```

Out[103]: {'rfc__criterion': 'gini',
           'rfc__max_depth': 9,
           'rfc__max_features': 9,
           'rfc__min_samples_leaf': 5,
           'rfc__n_estimators': 10}

```

```

In [104]: 1 gs_rfc.score(X_train, y_train)

```

```

Out[104]: 0.5083726150392817

```

### Evaluate gridsearch results

```

In [105]: 1 y_hat_gs_rfc = gs_rfc.predict(X_test)

```

In [106]: 1 `print(classification_report(y_test, y_hat_gs_rfc))`

	precision	recall	f1-score	support
functional	0.70	0.50	0.58	8098
functional needs repair	0.13	0.56	0.21	1074
non functional	0.62	0.48	0.54	5678
accuracy			0.50	14850
macro avg	0.48	0.51	0.44	14850
weighted avg	0.63	0.50	0.54	14850

In [561]: 1 `# Save the rfc gridsearch model model as a pickle in a file`  
 2 `joblib.dump(gs_rfc, 'gs_rfc.pkl')`

Out[561]: ['gs\_rfc.pkl']

### Results summary on Random Forest Gridsearch

Accuracy decreased significantly, perhaps as a result of using 10 `n_estimators` rather than the default 100 to cut down on processing time. This model though suggests where to explore for `max_depth`, and `samples leaf` and `split`.

### Use Randomized Search

```
In [137]: 1
2 from sklearn.model_selection import RandomizedSearchCV
3
4 # Based in previous gridsearch, optimize for max depth, min sample
5
6 random_grid = {
7     'rfc__bootstrap': [True],
8     'rfc__max_depth': [10, 20, 50, 100],
9     'rfc__max_features': ['auto', 'sqrt'],
10    'rfc__min_samples_leaf': [1, 2, 4],
11    'rfc__min_samples_split': [2, 5, 10],
12    'rfc__n_estimators': [10, 100]
13 }
14
15 random = RandomizedSearchCV(estimator = rfc_model_pipe,
16                             param_distributions = random_grid, n_jobs = -1,
17                             verbose = 2, random_state = 42, cv=3)
```

In [138]: 1 random.fit(X\_train, y\_train)

Fitting 3 folds for each of 10 candidates, totalling 30 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n\_jobs=-1)]: Done 30 out of 30 | elapsed: 124.2min finished

Out[138]: RandomizedSearchCV(cv=3,  
 estimator=Pipeline(steps=[('ct',  
 ColumnTransformer(remain  
 der='passthrough',  
 transformers=[('subpipe\_num',  
 Pipeline(steps=[('num\_impute',  
 SimpleImputer()),  
 ('ss',  
 StandardScaler())]),  
 [0,  
 2,  
 4,  
 5,  
 12]),  
 ('subpipe\_cat',  
 Pipeline(steps=[('cat\_impute',  
 SimpleImputer(fill\_value='missing',  
 strategy='constant'))],  
 ('ohe',  
 OneHotEncoder(handle\_unknow...  
 20,  
 21,  
 22,  
 23,  
 24,

```

25,
26]]]]),
                                ('sm', SMOTE(random_state=42)),
                                ('rfc',
                                 RandomForestClassifier
                                 (random_state=42))),
                                n_jobs=-1,
                                param_distributions={'rfc__bootstrap': [True],
                                                    'rfc__max_depth': [10, 20, 5
0, 100],
                                                    'rfc__max_features': ['auto',
'sqrt'],
                                                    'rfc__min_samples_leaf': [1,
2, 4],
                                                    'rfc__min_samples_split': [2,
5, 10],
                                                    'rfc__n_estimators': [10, 10
0]},
                                random_state=42, verbose=2)

```

```

In [140]: 1 # Best paramters from randomized search on RFC
          2 random.best_params_

```

```

Out[140]: {'rfc__n_estimators': 100,
           'rfc__min_samples_split': 5,
           'rfc__min_samples_leaf': 2,
           'rfc__max_features': 'auto',
           'rfc__max_depth': 100,
           'rfc__bootstrap': True}

```

### Evaluate randomized search best parameter results

```

In [139]: 1 random.score(X_train, y_train)

```

```

Out[139]: 0.8312457912457912

```

```

In [141]: 1 y_hat_random = random.predict(X_test)

```

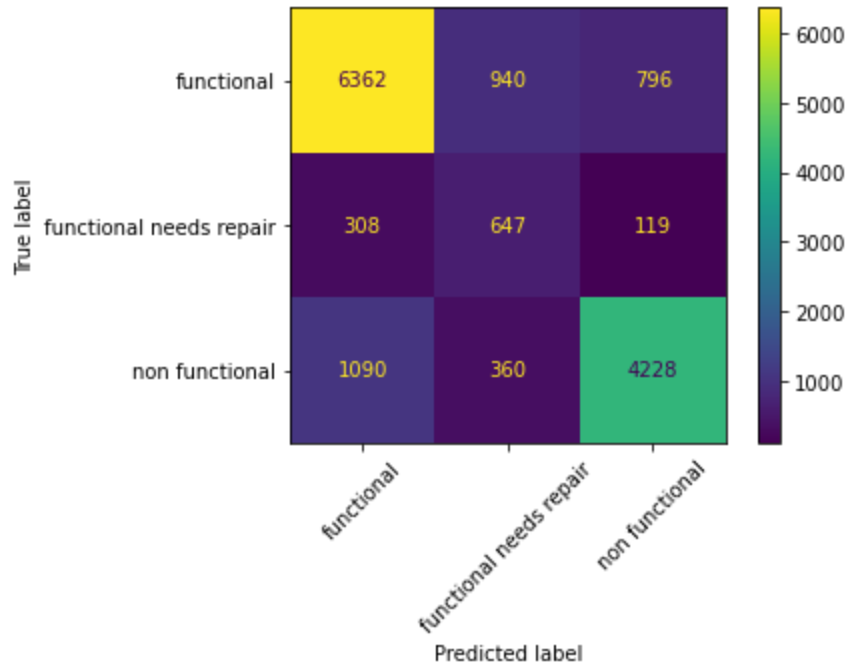
```

In [142]: 1 print(classification_report(y_test, y_hat_random))

```

	precision	recall	f1-score	support
functional	0.82	0.79	0.80	8098
functional needs repair	0.33	0.60	0.43	1074
non functional	0.82	0.74	0.78	5678
accuracy			0.76	14850
macro avg	0.66	0.71	0.67	14850
weighted avg	0.79	0.76	0.77	14850

```
In [569]: 1 plot_confusion_matrix(random, X_test, y_test, xticks_rotation = 45
```



```
In [562]: 1 # Save the model as a pickle in a file
          2 joblib.dump(random, 'random_rfc.pkl')
```

```
Out[562]: ['random_rfc.pkl']
```

### Gridsearch based on randomized results

```
In [571]: 1 # Based on randomized search conduct one more gridsearch
          2 params = {
          3     'rfc__n_estimators': [100],
          4     'rfc__min_samples_leaf': [2, 3],
          5     'rfc__max_depth': [100, 150],
          6     'rfc__min_samples_split': [3, 5, 7],
          7     'rfc__max_features': ['auto']
          8 }
          9
         10 gs_rfc_2 = GridSearchCV(estimator=rfc_model_pipe,
         11                          param_grid=params, n_jobs = -1,
         12                          verbose = 2, cv = 3)
         13
```

In [572]: 1 gs\_rfc\_2.fit(X\_train, y\_train)

Fitting 3 folds for each of 12 candidates, totalling 36 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n\_jobs=-1)]: Done 36 out of 36 | elapsed: 129.6min finished

Out [572]: GridSearchCV(cv=3,  
 estimator=Pipeline(steps=[('ct',  
 ColumnTransformer(remainder='transformer',  
 transformer  
 passthrough',  
 s=[('subpipe\_num',  
 Pipeline(steps=[('num\_impute',  
 SimpleImputer()),  
 ('ss',  
 StandardScaler())]),  
 [0, 2,  
 4, 5,  
 12]),  
 ('subpipe\_cat',  
 Pipeline(steps=[('cat\_impute',  
 SimpleImputer(fill\_value='missing',  
 strategy='constant'))],  
 ('ohe',  
 OneHotEncoder(handle\_unknown='ignore',  
 22,  
 23,  
 24,  
 25,  
 26])))),  
 ('sm',  
 SMOTE(n\_jobs=-1, random\_state=42)),  
 ('rfc',  
 RandomForestClassifier(max\_de

```

pth=100,
mples_leaf=2,
mples_split=3,
=-1,
_state=42)))),
    n_jobs=-1,
    param_grid={'rfc__max_depth': [100, 150],
                'rfc__max_features': ['auto'],
                'rfc__min_samples_leaf': [2, 3],
                'rfc__min_samples_split': [3, 5, 7],
                'rfc__n_estimators': [100]},
    verbose=2)

```

```

In [573]: 1 # Score the model on training data
          2 gs_rfc_2.score(X_train, y_train)

```

Out[573]: 0.8340291806958474

```

In [574]: 1 # examine best paramters
          2 gs_rfc_2.best_params_

```

Out[574]: {'rfc\_\_max\_depth': 100,  
 'rfc\_\_max\_features': 'auto',  
 'rfc\_\_min\_samples\_leaf': 2,  
 'rfc\_\_min\_samples\_split': 3,  
 'rfc\_\_n\_estimators': 100}

```

In [575]: 1 # create predicted target using test set
          2 y_hat_rfc_2 = gs_rfc_2.predict(X_test)

```

```

In [576]: 1 print(classification_report(y_test, y_hat_rfc_2))

```

	precision	recall	f1-score	support
functional	0.82	0.79	0.81	8098
functional needs repair	0.34	0.60	0.43	1074
non functional	0.83	0.75	0.79	5678
accuracy			0.76	14850
macro avg	0.66	0.71	0.67	14850
weighted avg	0.79	0.76	0.77	14850


```

In [577]: 1 # select as final model
          2 final_rfc_model = gs_rfc_2


```



```
In [787]: 1 # Plot confusion matrix with percentages
2 plot_confusion_matrix(final_rfc_model,
3                       X_test, y_test,
4                       xticks_rotation = 45,
5                       normalize = 'true'
6                       );
```



```
In [792]: 1 # Plot matrix with case numbers
2 plot_confusion_matrix(final_rfc_model,
3                       X_test, y_test,
4                       xticks_rotation = 45,
5
6                       );
7
8 plt.savefig('final RFC model matrix')
```



```
In [579]: 1 # Save the model as a pickle in a file
2 joblib.dump(final_rfc_model, 'final_rfc_model.pkl')
```

```
Out[579]: ['final_rfc_model.pkl']
```

## 5. Gradient Boost Model

Compare a default parameter gradient boost model against the RFC final model and select final model.

```
In [184]: 1 # Gradient Boost model
          2 gbc_model_pipe = Pipeline([('ct', CT), ('gbc', GradientBoostingCla
          3
          4 gbc_model_pipe.fit(X_train, y_train)
```

```
Out[184]: Pipeline(steps=[('ct',
                           ColumnTransformer(remainder='passthrough',
                                                transformers=[('subpipe_num',
                                                                Pipeline(steps=[('n
um_impute',
                                                                Si
mpleImputer()),
                                                                ('s
s',
                                                                St
andardScaler())])),
                           [0, 2, 4, 5, 12]),
                           ('subpipe_cat',
                            Pipeline(steps=[('c
at_impute',
                            Si
mpleImputer(fill_value='missing',
strategy='constant'))],
                           ('o
neHotEncoder(handle_unknown='ignore',
sparse=False))])),
                           [1, 3, 6, 7, 8, 9,
10, 11, 13,
14, 15, 16, 17, 1
8, 19, 20,
21, 22, 23, 24, 2
5, 26])])),
                           ('gbc', GradientBoostingClassifier(random_state=4
2)))]
```

```
In [185]: 1 gbc_model_pipe.score(X_train, y_train)
```

```
Out[185]: 0.7612570145903479
```


```
In [186]: 1 y_hat_gbc = gbc_model_pipe.predict(X_test)
          2
          3 print(classification_report(y_test, y_hat_gbc))
```

	precision	recall	f1-score	support
functional	0.72	0.93	0.81	8098
functional needs repair	0.66	0.17	0.27	1074
non functional	0.85	0.61	0.71	5678
accuracy			0.75	14850
macro avg	0.74	0.57	0.60	14850
weighted avg	0.76	0.75	0.73	14850

```
In [193]: 1 gbc_model_pipe.named_steps['gbc']
```

```
Out[193]: GradientBoostingClassifier(random_state=42)
```

```
In [568]: 1 plot_confusion_matrix(gbc_model_pipe, X_test, y_test, xticks_rotat
```



```
In [565]: 1 # Save the model as a pickle in a file
          2 joblib.dump(gbc_model_pipe, 'gbc_model.pkl')
```

```
Out[565]: ['gbc_model.pkl']
```

## 6. Data Visualizations

Create three data visualizations to communicate findings to Water Aid.

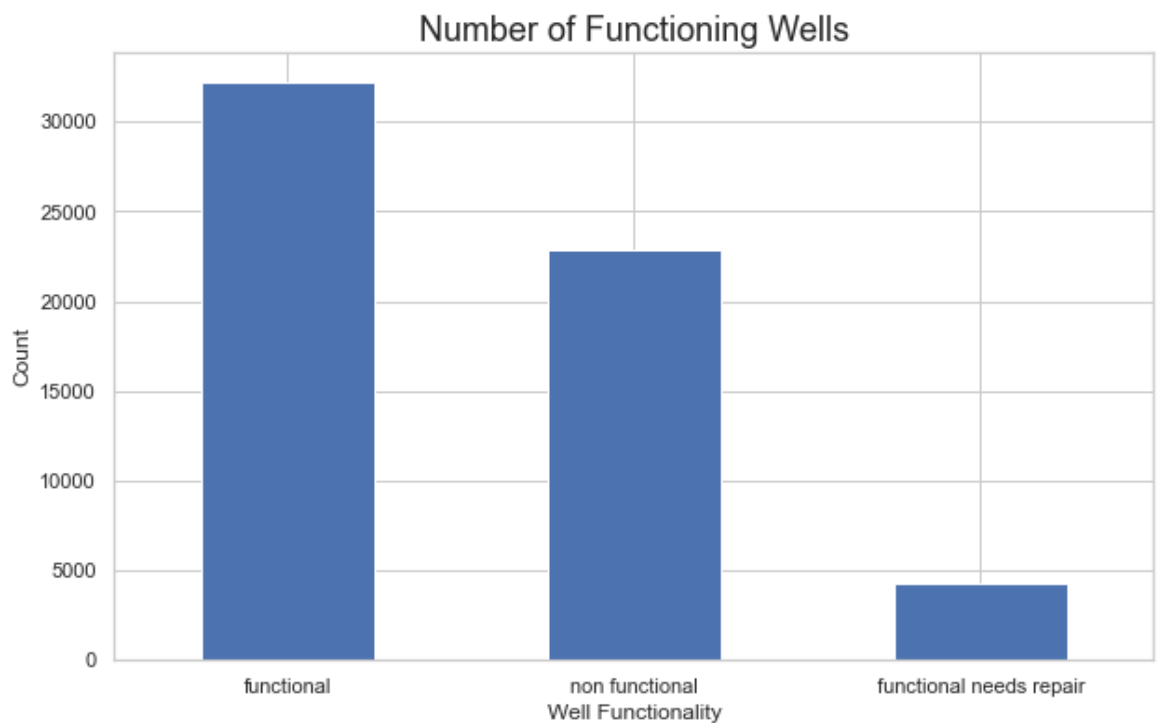
```
In [797]: 1 # Examine full dataframe columns
          2 df.columns
```

```
Out[797]: Index(['id', 'amount_tsh', 'date_recorded', 'funder', 'gps_height',
                'installer', 'longitude', 'latitude', 'wpt_name', 'num_private',
                'basin', 'subvillage', 'region', 'region_code', 'district_code',
                'lga', 'ward', 'population', 'public_meeting', 'recorded_by',
                'scheme_management', 'scheme_name', 'permit', 'construction_year',
                'extraction_type', 'extraction_type_group', 'extraction_type_class',
                'management', 'management_group', 'payment', 'payment_type',
                'water_quality', 'quality_group', 'quantity', 'quantity_group',
                'source', 'source_type', 'source_class', 'waterpoint_type',
                'waterpoint_type_group', 'status_group'],
                dtype='object')
```

```

In [816]: 1 # Visualize well status count from dataset
          2
          3 df.status_group.value_counts().plot(kind="bar")
          4 plt.title("Number of Functioning Wells", fontsize= 18)
          5 plt.xlabel("Well Functionality", fontsize = 12)
          6 plt.xticks(rotation=0)
          7 plt.ylabel("Count", fontsize = 12)
          8
          9 plt.savefig('Number of Functioning Wells.png')
         10
         11 plt.show();
         12
         13

```



```

In [799]: 1 # Plot four shared cross country water basins in the region
          2 # Plot visual showing basins and functional wells
          3 df_basin = df[df['basin'].isin(['Lake Nyasa', 'Lake Victoria', 'La
          4                                         'Ruvuma / Southern Coast',
          5                                         ])]

```

```

In [800]: 1 df_basin.basin.value_counts()

```

```


Out[800]: Lake Victoria          10248
          Lake Tanganyika         6432
          Lake Nyasa              5085
          Ruvuma / Southern Coast  4493
          Name: basin, dtype: int64

```

```
In [801]: 1 df_basin.status_group.value_counts()
```

```
Out[801]: functional          13201  
non functional             10750  
functional needs repair     2307  
Name: status_group, dtype: int64
```

```
In [829]: 1 # Use Seaborn to and stacked histogram to show the four basins and  
2  
3 sns.set_theme()  
4  
5 sns.set(rc={"figure.figsize":(10, 6)})  
6 sns.set_style('whitegrid')  
7  
8 sns.histplot(data = df_basin, x = 'basin', hue = 'status_group',  
9             bins = 10, binwidth = 6, palette = 'GnBu', legend = '  
10             multiple = 'stack')  
11  
12  
13 plt.title("Functioning Wells by Basin", fontsize= 24)  
14 plt.xlabel(None)  
15 plt.ylabel("Count", fontsize = 16)  
16 plt.xticks(rotation = 0, fontsize = 12)  
17  
18 plt.savefig('functioning wells by basin.png')  
19  
20 plt.show();  
21  
22
```



```
In [803]: 1 # Examine well status for two basin recommendations: Lake Victoria
          2 df_victoria = df[df['basin'].isin(['Lake Victoria'])]
          3
          4 df_victoria.status_group.value_counts()
```

```
Out[803]: functional          5100
          non functional      4159
          functional needs repair    989
          Name: status_group, dtype: int64
```

```
In [804]: 1 # Create dataframe for Ruvuma Basin
          2 df_ruvuma = df[df['basin'].isin(['Ruvuma / Southern Coast'])]
          3
          4 #Examine well function status for Ruvuma Basin
          5 print(df_ruvuma.status_group.value_counts(normalize = True))
          6 print()
          7 print(df_ruvuma.status_group.value_counts())
```

```
non functional          0.555753
functional              0.371689
functional needs repair  0.072557
Name: status_group, dtype: float64
```

```
non functional          2497
functional              1670
functional needs repair   326
Name: status_group, dtype: int64
```

```
In [805]: 1 df_ruvuma.head()
```

```
Out[805]:
```

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	v
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	
26	55012	500.0	2013-01-16	Sobodo	200	Kilolo Star	39.370777	-9.942532	
46	45111	20.0	2013-02-05	Lga	240	LGA	39.087415	-11.000604	M
91	62591	0.0	2013-01-20	Jica	212	Kokeni	38.962945	-10.476566	
98	33379	0.0	2013-02-19	Danida	1000	DWE	35.542173	-10.808853	

5 rows × 41 columns



```
In [806]: 1 #examine descriptive statistics for Ruvuma
          2 df_ruvuma.describe()
```

Out[806]:


	id	amount_tsh	gps_height	longitude	latitude	num_private	region
<b>count</b>	4493.000000	4493.000000	4493.000000	4493.000000	4493.000000	4493.000000	4493.0
<b>mean</b>	37322.257067	228.390385	410.640329	38.316789	-10.485215	0.124861	52.0
<b>std</b>	21489.456338	777.985990	338.566284	1.549237	0.591604	6.745120	38.0
<b>min</b>	19.000000	0.000000	-90.000000	34.889771	-11.649440	0.000000	8.0
<b>25%</b>	18908.000000	0.000000	164.000000	37.244214	-10.850966	0.000000	10.0
<b>50%</b>	37228.000000	0.000000	342.000000	38.935668	-10.626269	0.000000	80.0
<b>75%</b>	55874.000000	50.000000	585.000000	39.448147	-10.250908	0.000000	90.0
<b>max</b>	74247.000000	15000.000000	1641.000000	40.345193	-8.496806	450.000000	99.0

```
In [807]: 1 # Build dataframe ruvuma basin with construction year
          2 df_ruvuma_built = df_ruvuma[df_ruvuma.construction_year != 0]
```

```

In [818]: 1 # Create data visualization, histogram, for functioning wells in R
          2 sns.set_theme()
          3
          4 sns.set(rc={"figure.figsize":(10, 6)})
          5 sns.set_style('whitegrid')
          6
          7 sns.histplot(data = df_ruvuma_built, x = 'construction_year', hue
          8                 bins = 20, binwidth = 5, palette = 'GnBu', legend = '
          9                 multiple = 'stack')
         10
         11
         12 plt.title("Functioning Wells by Year Constructed", fontsize= 18)
         13 plt.xlabel('Ruvuma Basin', fontsize = 16)
         14 plt.xticks(rotation=0)
         15 plt.ylabel("Count", fontsize = 16)
         16 plt.xticks(rotation = 45)
         17
         18 plt.savefig('functioning wells ruvuma.png')
         19
         20 plt.show()
         21
         22 ;

```



Out [818]: ..

```

In [809]: 1 # Dataframe for older wells in Ruvuma, 1975 to 1990
          2 df_ruvuma_old_wells = df_ruvuma.loc[(df_ruvuma.construction_year >
          3                 (df_ruvuma.construction_year <=

```

```
In [810]: 1 # Percentage of wells in need of repair
          2 df_ruvuma_old_wells.status_group.value_counts(normalize = True)

Out[810]: non functional      0.699024
          functional         0.236646
          functional needs repair 0.064331
          Name: status_group, dtype: float64
```

## Final Summary

The initial Logistic Regression model with default parameters delivered the following scores:

		precision	recall	f1-score	sup
port	functional	0.78	0.88	0.83	80
	functional needs repair	0.55	0.25	0.34	10
	non functional	0.80	0.74	0.77	56

The final Random Forest Classifier model with class imbalance adjustments and hyperparameter tuning delivered the following scores:

		precision	recall	f1-score	sup
port	functional	0.82	0.79	0.81	80
	functional needs repair	0.34	0.60	0.43	10
	non functional	0.83	0.75	0.79	56

The Random Forest Classifier was trained using both Randomized Search and Grid Search. Here are the final hyperparameter adjustments:

`n_estimators = 100, (default)`

`max_depth = 100,`

`max_features = 'auto',`

`min_samples_leaf = 2,`

`min_samples_split = 3`

The precision scores increased by 4% for functional wells and increased by 3% for non-functional wells. The precision score fell quite dramatically by 21% for wells in need of repair.

This may be due to wells in need of repair often being classified as functioning wells, their features appear to look much like functioning wells. Despite this the recall score for wells in need of repair improved dramatically, by 35%.

The improved model delivers trade offs. As precision scores increase for functional and non-functional wells, precision for wells in need of repair decreases. But, it needs to be noted, recall for wells in need of repair increases.

Water Aid's use of the model will be primarily for precision - true positive identification for functioning and non-functioning wells. They can still make use of the model for wells in need of repair, but the results for that class need to be understood in terms of recall - the ability of a model to find all the relevant cases within a data set.

## Recommendations

### 1. Model Use

Wells in shared basins should make use of the model's precision metric to accurately identify non-functioning wells. Chances are 83% that they will be right which will help in making use of programming resources to repair the wells.

### 2. Basin Location

The model makes significant use of data in Tanzanian water basins that also span nearby countries. The Ruvuma Basin stretches from southern Tanzania to northern Mozambique. The Ruvuma basin contains more than 55% nonfunctioning wells and over 7% of wells need repair. Considering Water Aid has a presence in both countries, and efforts by other NGOs and governments of both countries to manage this space along transboundary lines, the Ruvuma basin offers an opportunity to make an impact using this model.

### 3. Well Age

In the Ruvuma Basin the proportion of older wells in need of repair compared to newer wells in need of repair is much higher. Wells built between 1975 to 1990 should be targeted first. These older non-functional wells account for nearly 70% of all wells built during this time frame.